

Asked an AI for ideas.

Got an answer that felt
like someone else's.

Every new session starts from zero.

No memory of who you are,
what you work on, or how you
write.

It fills that gap with a guess.

**Andrej Karpathy,
ex-OpenAI, ex-Tesla,
published a fix for
this as a free
public gist.**

He calls it the LLM wiki.

The folders:

/ raw - unsorted stuff you drop in, notes, clips, transcripts.

/ wiki - short AI-written pages, one per topic, under about 500 words.

/ journal - dated entries, like a diary.

/ crm - contacts and whos who.

agents.md, index.md, log.md - three plain files that run the whole system.

The loop:

1. you drop raw material into /raw
2. on a schedule, the AI reads it, pulls out the ideas, writes short wiki pages, links them together
3. it updates index.md and log.md so nothing gets lost
4. the original file moves to /raw/processed – it never repeats work.

**An AI reasons
better over one
tidy 500-word page
than a raw 8,000-word
transcript.**

So instead of you re-explaining
yourself every session, the AI
keeps its own compressed notes
and reads those instead.

To start one yourself:

1. make the five folders above.
2. write agents.md in plain English, how to file new material, how to answer questions, how to update the index.
3. start dropping things into /raw and let the AI run the loop.

full spec: gist.github.com/karpathy, search LLM wiki.

**The fix for
'it doesn't
know me'
turned out to be:
give it a notebook,
and tell it how
to use it.**